

# Dissert-final-exp

Gabriel Robaina

2024-03-21


## Dissert experiments

### Write operations on large files


```
## [1] "ECDF of write latency for a 1 GB file"
```

```
## 'summarise()' has grouped output by 'system_name'. You can override using the  
## '.groups' argument.
```

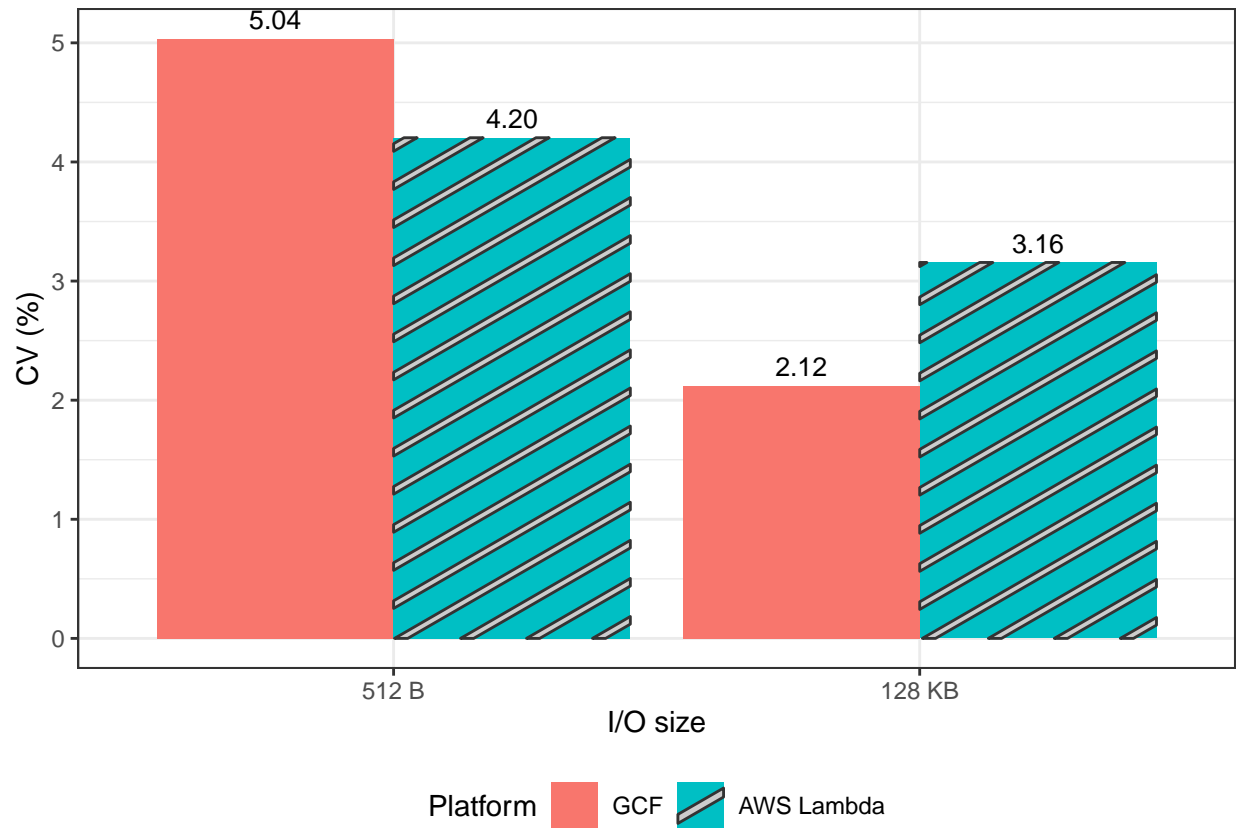
```
## # A tibble: 4 x 3  
## # Groups:   system_name [2]  
##   system_name io_size_bytes median_latency  
##   <chr>         <int64>         <dbl>  
## 1 GCF_DD           512             9.54  
## 2 GCF_DD        128000             4.07  
## 3 LAMBDA_DD         512             7.85  
## 4 LAMBDA_DD        128000             3.60
```



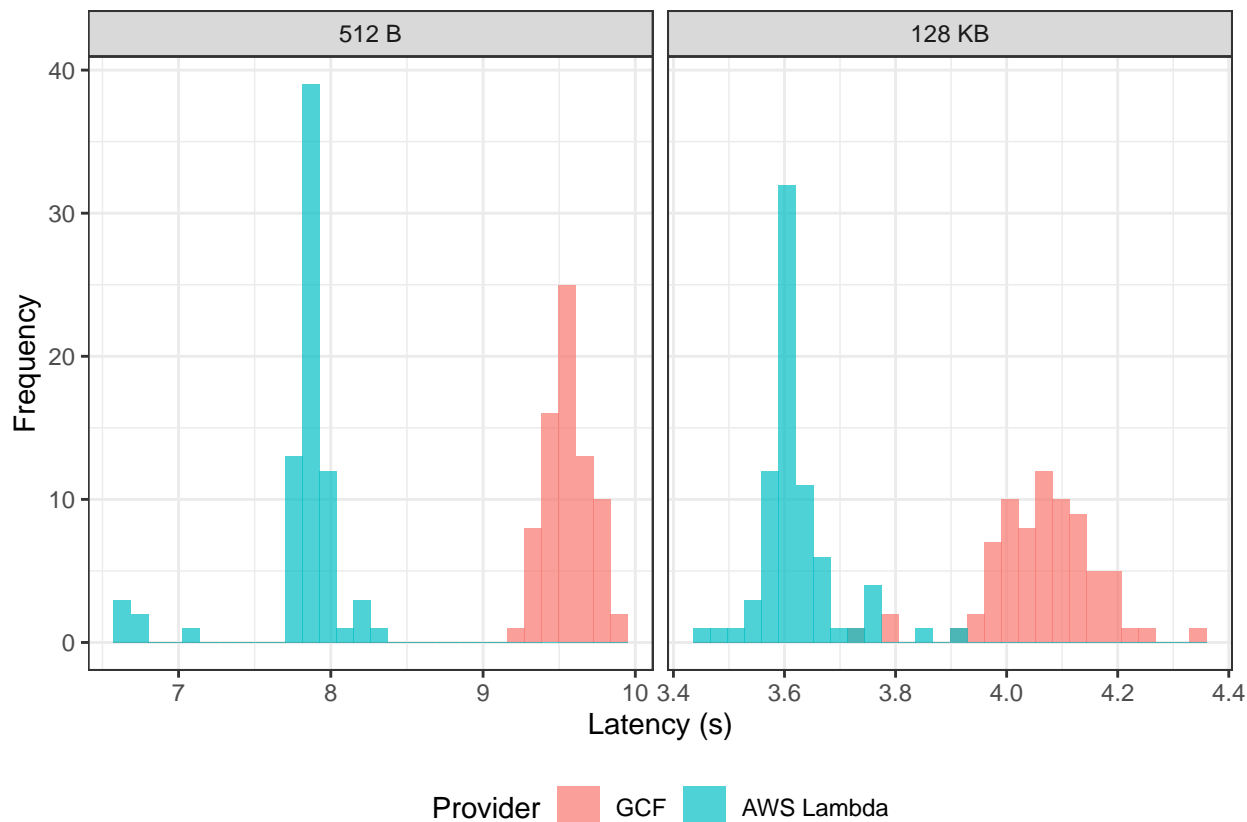
##	provider	io_size	lower	upper	mean
## 1	LAMBDA_DD	512	7.716750	7.868251	7.792501
## 2	LAMBDA_DD	128000	3.604048	3.635907	3.619977
## 3	GCF_DD	512	9.520569	9.586562	9.553566
## 4	GCF_DD	128000	4.041809	4.086982	4.064395



```
## [1] "CV for writes"
```



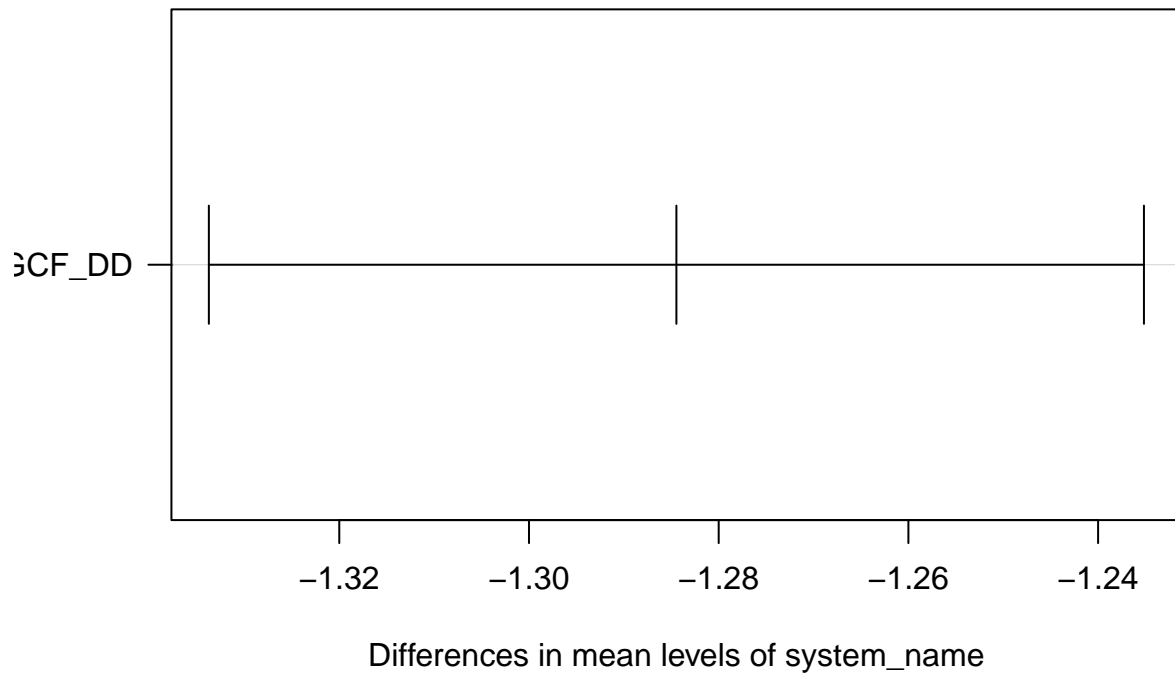
```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```




```
## [1] "1 GB file for WRITE operations"
##               Df Sum Sq Mean Sq F value Pr(>F)
## system_name      1    247      247    2619 <2e-16 ***
## io_size_bytes_fac 1   3790     3790  40109 <2e-16 ***
## system_name:io_size_bytes_fac 1    109      109    1152 <2e-16 ***
## Residuals      596      56        0
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## Tukey multiple comparisons of means
## 95% family-wise confidence level
##
## Fit: aov(formula = latency_seconds ~ system_name * io_size_bytes_fac, data = res)
##
## $system_name
##               diff          lwr          upr p adj
## LAMBDA_DD-GCF_DD -1.284458 -1.333751 -1.235165      0
##
## $io_size_bytes_fac
##               diff          lwr          upr p adj
## 128000-512 -5.026602 -5.075895 -4.97731      0
##
## $'system_name:io_size_bytes_fac'
##               diff          lwr          upr p adj
## LAMBDA_DD:512-GCF_DD:512 -2.1361541 -2.2275998 -2.0447085      0
## GCF_DD:128000-GCF_DD:512 -5.8782985 -5.9697442 -5.7868529      0
## LAMBDA_DD:128000-GCF_DD:512 -6.3110603 -6.4025059 -6.2196146      0
```

##	GCF_DD:128000-LAMBDA_DD:512	-3.7421444	-3.8335900	-3.6506988	0
##	LAMBDA_DD:128000-LAMBDA_DD:512	-4.1749061	-4.2663518	-4.0834605	0
##	LAMBDA_DD:128000-GCF_DD:128000	-0.4327617	-0.5242074	-0.3413161	0


### 95% family-wise confidence level



### 95% family-wise confidence level



## 95% family-wise confidence level



Differences in mean levels of system\_name:io\_size\_bytes\_fac

```
## [1] "Sum of squares"
##                               Sum Sq
## system_name                  5.89
## io_size_bytes_fac            90.18
## system_name:io_size_bytes_fac 2.59
## Residuals                    1.34
## [1] "F_crit calculation"
##   factor_df residual_df   F_crit
## 1         1          596 3.857108
```


## Write operations on small files

```
## [1] "ECDF of write latency for a 10 KB file and 512 B I/O size"
```


```
## 'summarise()' has grouped output by 'system_name'. You can override using the
## '.groups' argument.
```

```
## # A tibble: 4 x 3
## # Groups:   system_name [2]
##   system_name resource_tier median_latency
##   <chr>         <chr>         <dbl>
## 1 GCF_DD       TIER_1             0.274
## 2 GCF_DD       TIER_5             0.276
## 3 LAMBDA_DD    TIER_1             0.264
## 4 LAMBDA_DD    TIER_5             0.265
```

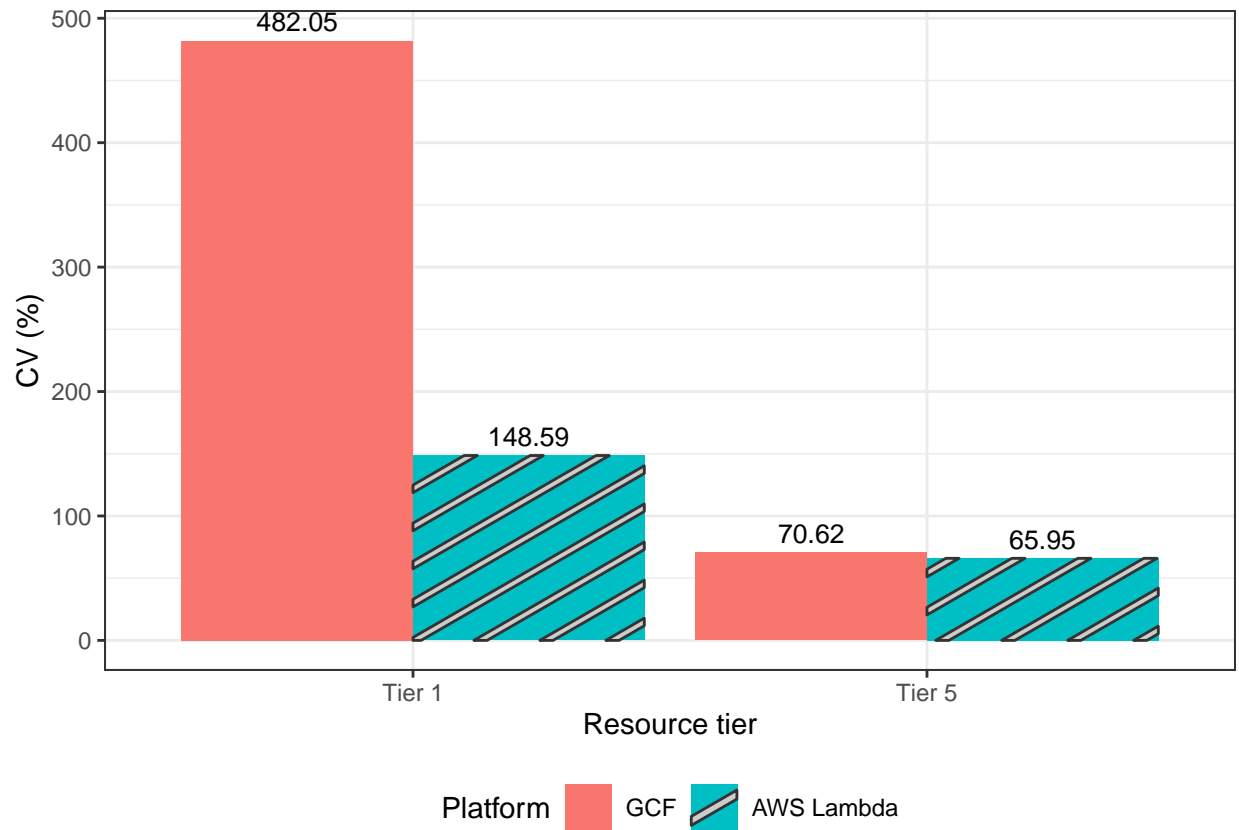





##	provider	resource_tier	lower	upper	mean
## 1	LAMBDA_DD	TIER_1	6.0075037	9.7961710	7.9018374
## 2	LAMBDA_DD	TIER_5	0.2547181	0.3153782	0.2850482
## 3	GCF_DD	TIER_1	0.9390005	7.5109078	4.2249541
## 4	GCF_DD	TIER_5	0.2739734	0.3444359	0.3092047



```
## [1] "CV for writes"
```



```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```




```
## [1] "10 KB file for WRITE operations"
##               Df Sum Sq Mean Sq F value    Pr(>F)
## system_name      1 0.00050 0.000500    3.621   0.0575 .
## resource_tier     1 0.00499 0.004987   36.093 3.27e-09 ***
## system_name:resource_tier 1 0.00051 0.000514    3.717   0.0543 .
## Residuals       596 0.08236 0.000138
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## Tukey multiple comparisons of means
## 95% family-wise confidence level
##
## Fit: aov(formula = latency_seconds ~ system_name * resource_tier, data = res)
##
## $system_name
##               diff               lwr               upr             p adj
## LAMBDA_DD-GCF_DD 0.001826363 -5.864643e-05 0.003711373 0.0575418
##
## $resource_tier
##               diff               lwr               upr             p adj
## TIER_5-TIER_1 -0.005766269 -0.007651279 -0.00388126      0
##
## $'system_name:resource_tier'
##               diff               lwr               upr
## LAMBDA_DD:TIER_1-GCF_DD:TIER_1 3.676883e-03 0.0001799041 0.0071738623
## GCF_DD:TIER_5-GCF_DD:TIER_1 -3.915749e-03 -0.0074127285 -0.0004187704
## LAMBDA_DD:TIER_5-GCF_DD:TIER_1 -3.939906e-03 -0.0074368851 -0.0004429269
```

```


## GCF_DD:TIER_5-LAMBDA_DD:TIER_1    -7.592633e-03 -0.0110896118 -0.0040956536
## LAMBDA_DD:TIER_5-LAMBDA_DD:TIER_1  -7.616789e-03 -0.0111137683 -0.0041198101
## LAMBDA_DD:TIER_5-GCF_DD:TIER_5     -2.415652e-05 -0.0035211356  0.0034728226
##                                     p adj
## LAMBDA_DD:TIER_1-GCF_DD:TIER_1     0.0349446
## GCF_DD:TIER_5-GCF_DD:TIER_1        0.0210927
## LAMBDA_DD:TIER_5-GCF_DD:TIER_1     0.0200063
## GCF_DD:TIER_5-LAMBDA_DD:TIER_1     0.0000002
## LAMBDA_DD:TIER_5-LAMBDA_DD:TIER_1  0.0000002
## LAMBDA_DD:TIER_5-GCF_DD:TIER_5     0.9999980

```


### 95% family-wise confidence level



### 95% family-wise confidence level



## 95% family-wise confidence level



Differences in mean levels of system\_name:resource\_tier


```
## [1] "Sum of squares"
##                               Sum Sq
## system_name                 0.57
## resource_tier                5.64
## system_name:resource_tier    0.58
## Residuals                   93.21
## [1] "F_crit calculation"
##   factor_df residual_df  F_crit
## 1         1          596 3.857108
```

## Read operations on large files

```
## [1] "ECDF of read latency for a 1 GB file"


## 'summarise()' has grouped output by 'system_name'. You can override using the
## '.groups' argument.

## # A tibble: 4 x 3
## # Groups:   system_name [2]
##   system_name io_size_bytes median_latency
##   <chr>      <int64>      <dbl>
## 1 GCF_DD        512         4.79
## 2 GCF_DD    128000         0.0815
## 3 LAMBDA_DD     512        265.
## 4 LAMBDA_DD    128000        29.8
```

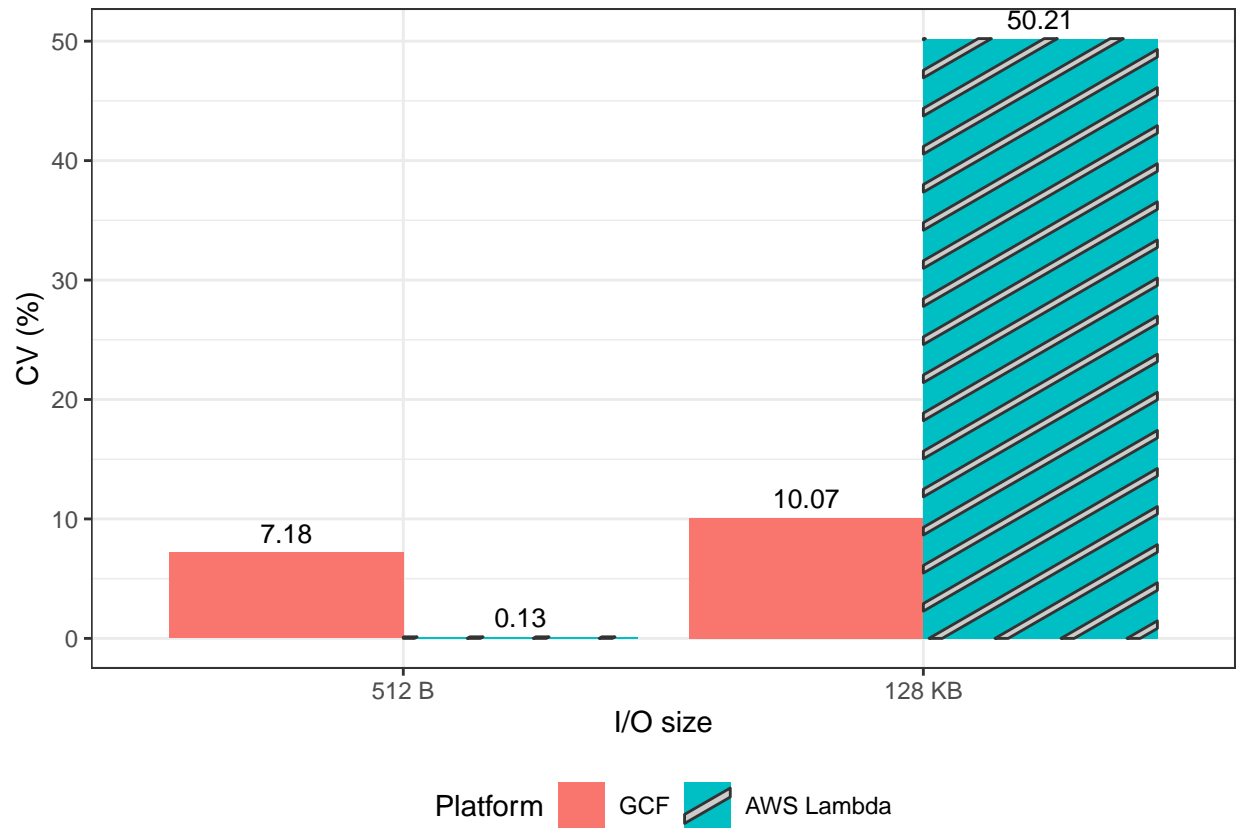


##	provider	io_size	lower	upper	mean
## 1	LAMBDA_DD	512	265.16117135	265.34426865	265.25272000
## 2	LAMBDA_DD	128000	20.82039683	25.96912930	23.39476307
## 3	GCF_DD	512	4.77805686	4.81451941	4.79628813
## 4	GCF_DD	128000	0.08060797	0.08479676	0.08270236

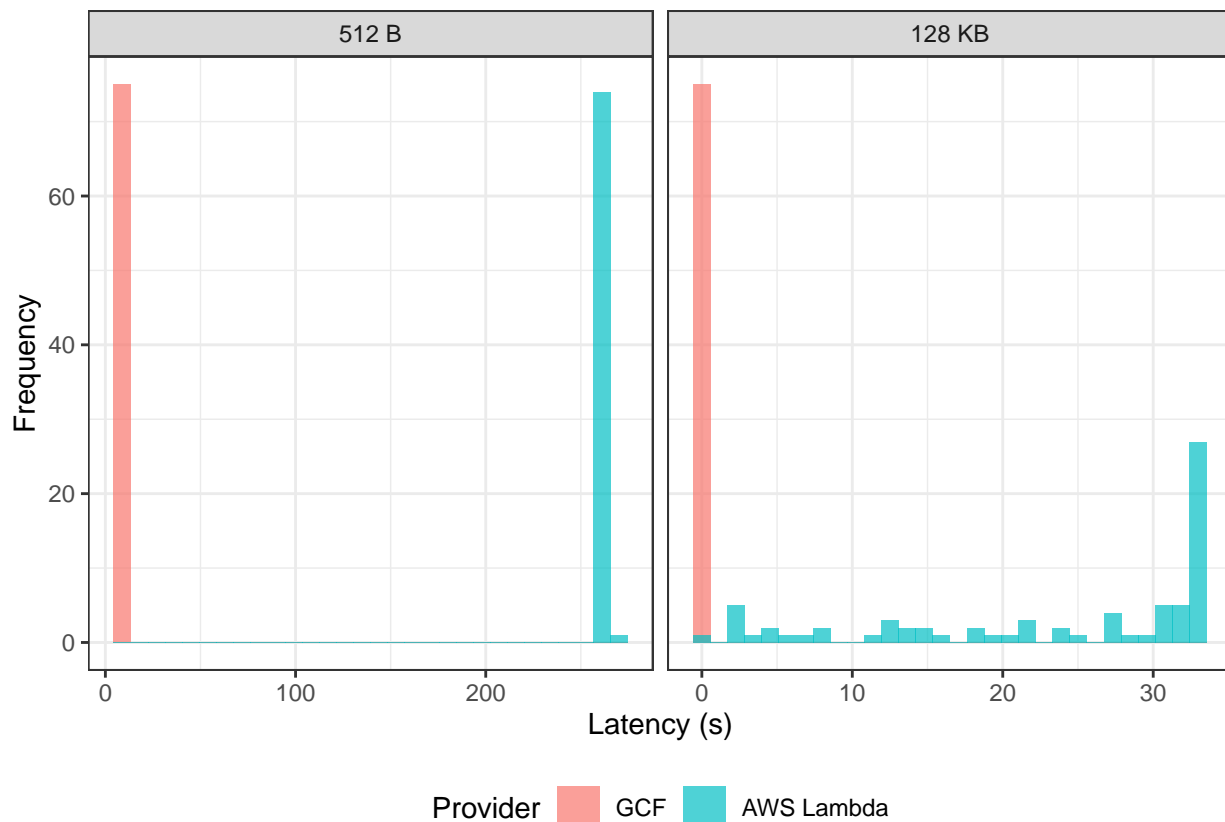




```
## [1] "CV for reads"
```




```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```




```
## [1] "1 GB file for READ operations"
##               Df Sum Sq Mean Sq F value Pr(>F)
## system_name    1 2982180 2982180   98403 <2e-16 ***
## io_size_bytes_fac    1 2313112 2313112   76326 <2e-16 ***
## system_name:io_size_bytes_fac    1 2130624 2130624   70304 <2e-16 ***
## Residuals      596   18062     30
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## Tukey multiple comparisons of means
## 95% family-wise confidence level
##
## Fit: aov(formula = latency_seconds ~ system_name * io_size_bytes_fac, data = res)
##
## $system_name
##              diff          lwr          upr p adj
## LAMBDA_DD-GCF_DD 141.0007 140.1179 141.8835     0
##
## $io_size_bytes_fac
##              diff          lwr          upr p adj
## 128000-512 -124.1803 -125.0631 -123.2975     0
##
## $'system_name:io_size_bytes_fac'
##              diff          lwr          upr p adj
## LAMBDA_DD:512-GCF_DD:512    260.181915 258.544244 261.819587     0
## GCF_DD:128000-GCF_DD:512    -4.999086  -6.636758 -3.361414     0
## LAMBDA_DD:128000-GCF_DD:512    16.820425  15.182754  18.458097     0
```

##	GCF_DD:128000-LAMBDA_DD:512	-265.181001	-266.818673	-263.543330	0
##	LAMBDA_DD:128000-LAMBDA_DD:512	-243.361490	-244.999162	-241.723819	0
##	LAMBDA_DD:128000-GCF_DD:128000	21.819511	20.181839	23.457183	0


### 95% family-wise confidence level



### 95% family-wise confidence level



## 95% family-wise confidence level




```
## [1] "Sum of squares"
##                               Sum Sq
## system_name                 40.06
## io_size_bytes_fac           31.07
## system_name:io_size_bytes_fac 28.62
## Residuals                    0.24
## [1] "F_crit calculation"
##   factor_df residual_df   F_crit
## 1         1           596 3.857108
```

## Read operations on small files

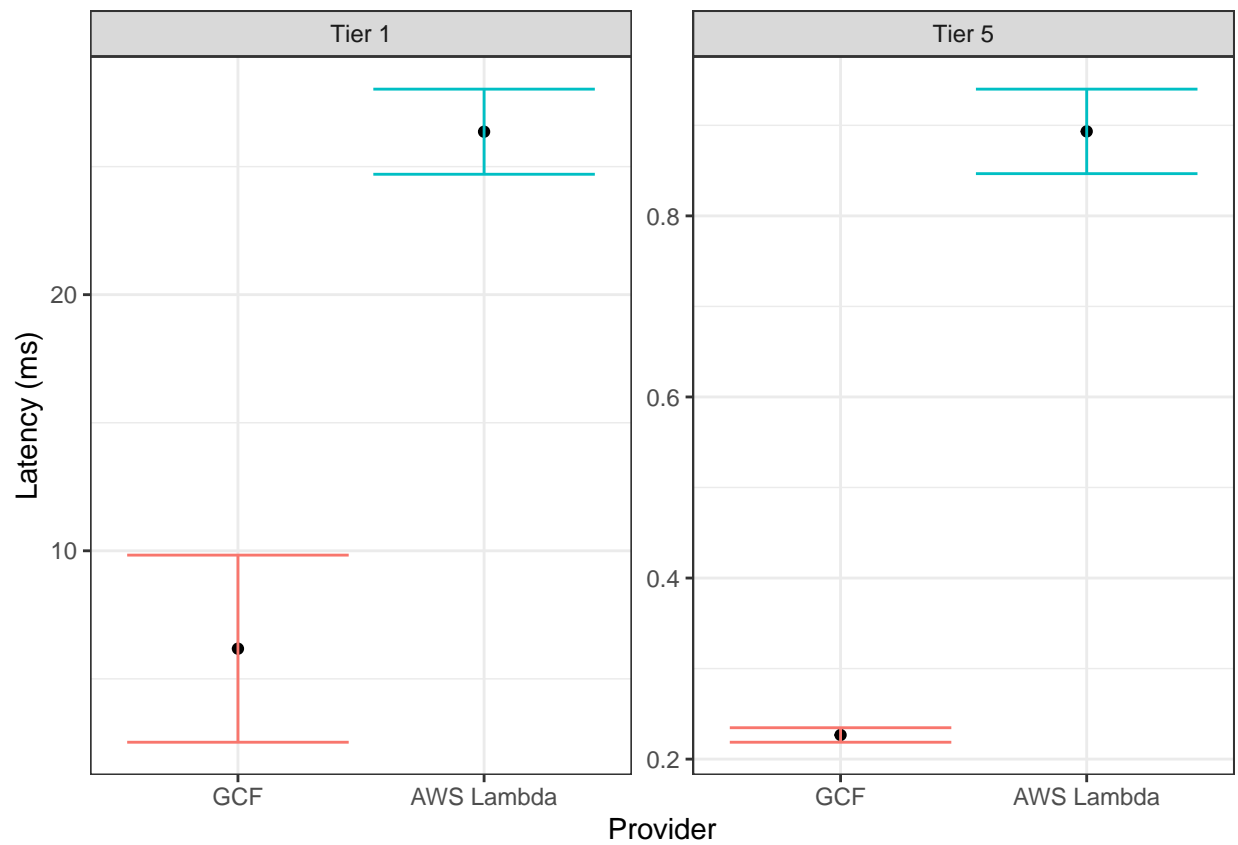
```
## [1] "ECDF of read latency for a 10 KB file and 512 B I/O size"
```

```
## 'summarise()' has grouped output by 'system_name'. You can override using the
## '.groups' argument.
```

```
## # A tibble: 4 x 3
## # Groups:   system_name [2]
##   system_name resource_tier median_latency
##   <chr>         <chr>         <dbl>
## 1 GCF_DD       TIER_1             0.211
## 2 GCF_DD       TIER_5             0.214
## 3 LAMBDA_DD    TIER_1            20.3
## 4 LAMBDA_DD    TIER_5             0.839
```

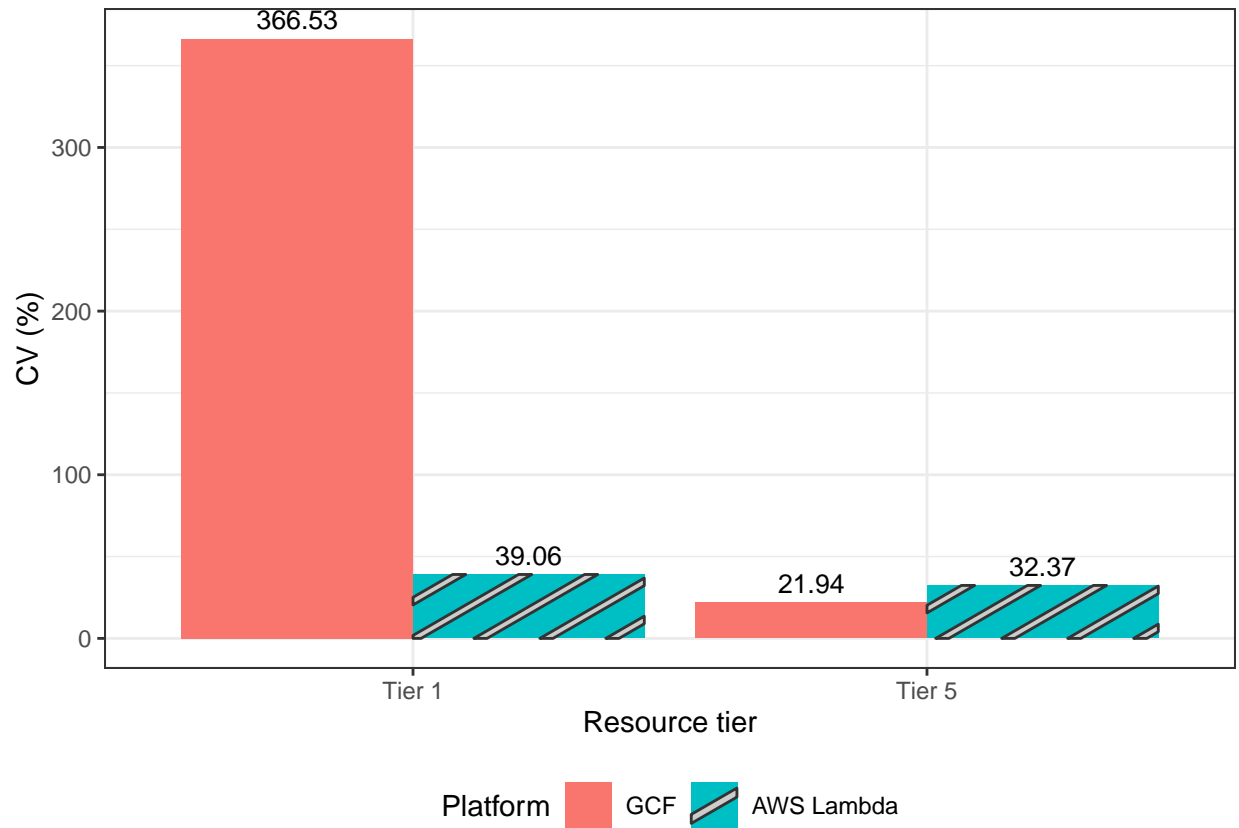


##	provider	resource_tier	lower	upper	mean
## 1	LAMBDA_DD	TIER_1	24.7018439	28.0245315	26.3631877
## 2	LAMBDA_DD	TIER_5	0.8466467	0.9399647	0.8933057
## 3	GCF_DD	TIER_1	2.5244415	9.8312064	6.1778239
## 4	GCF_DD	TIER_5	0.2186401	0.2346863	0.2266632




## [1] "CV for reads"






```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



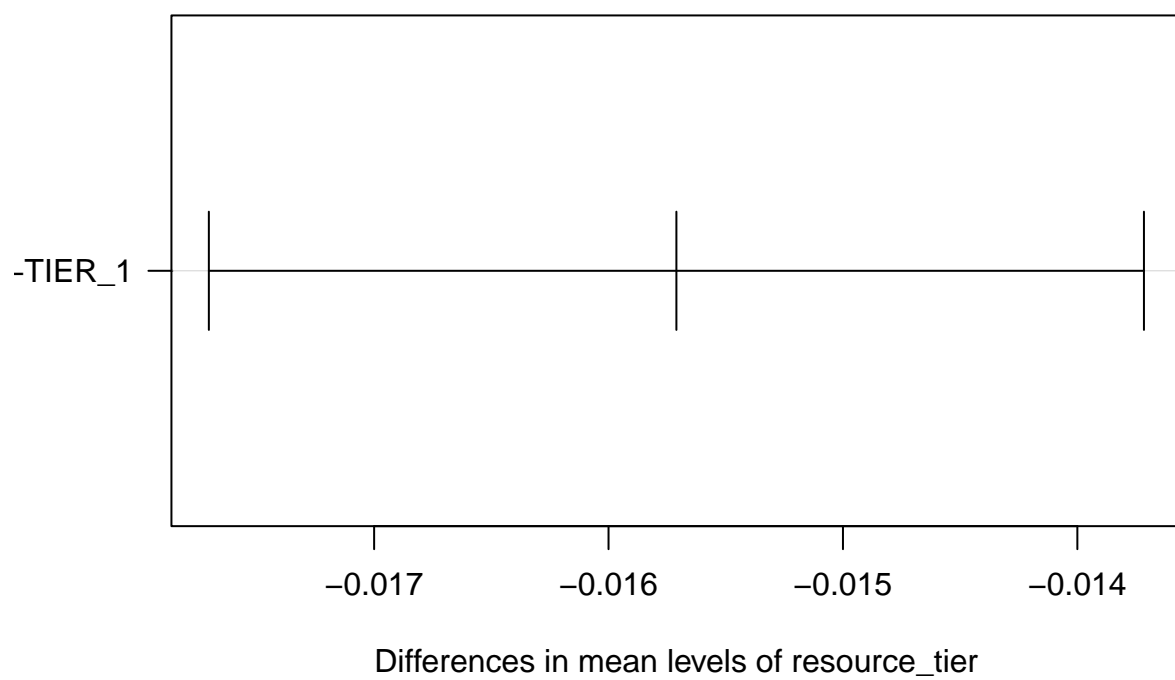
```
## [1] "10 KB file for READ operations"
##               Df Sum Sq Mean Sq F value Pr(>F)
## system_name      1 0.01631 0.01631   105.39 <2e-16 ***
## resource_tier     1 0.03702 0.03702   239.30 <2e-16 ***
## system_name:resource_tier 1 0.01429 0.01429   92.34 <2e-16 ***
## Residuals       596 0.09221 0.00015
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## Tukey multiple comparisons of means
## 95% family-wise confidence level
##
## Fit: aov(formula = latency_seconds ~ system_name * resource_tier, data = res)
##
## $system_name
##               diff              lwr              upr p adj
## LAMBDA_DD-GCF_DD 0.010426 0.00843142 0.01242059      0
##
## $resource_tier
##               diff              lwr              upr p adj
## TIER_5-TIER_1 -0.01571052 -0.0177051 -0.01371594      0
##
## $'system_name:resource_tier'
##               diff              lwr              upr
## LAMBDA_DD:TIER_1-GCF_DD:TIER_1      0.0201853638 0.016485109 0.023885619
## GCF_DD:TIER_5-GCF_DD:TIER_1      -0.0059511607 -0.009651416 -0.002250906
## LAMBDA_DD:TIER_5-GCF_DD:TIER_1      -0.0052845182 -0.008984773 -0.001584263
```

```
## GCF_DD:TIER_5-LAMBDA_DD:TIER_1    -0.0261365245 -0.029836780 -0.022436269
## LAMBDA_DD:TIER_5-LAMBDA_DD:TIER_1 -0.0254698820 -0.029170137 -0.021769627
## LAMBDA_DD:TIER_5-GCF_DD:TIER_5      0.0006666425 -0.003033613  0.004366898
##                                     p adj
## LAMBDA_DD:TIER_1-GCF_DD:TIER_1      0.00000000
## GCF_DD:TIER_5-GCF_DD:TIER_1          0.0002287
## LAMBDA_DD:TIER_5-GCF_DD:TIER_1       0.0014521
## GCF_DD:TIER_5-LAMBDA_DD:TIER_1       0.0000000
## LAMBDA_DD:TIER_5-LAMBDA_DD:TIER_1    0.0000000
## LAMBDA_DD:TIER_5-GCF_DD:TIER_5       0.9668203
```


### 95% family-wise confidence level



**95% family-wise confidence level**



## 95% family-wise confidence level



```
## [1] "Sum of squares"
##                               Sum Sq
## system_name                  10.20
## resource_tier                 23.16
## system_name:resource_tier     8.94
## Residuals                     57.69
## [1] "F_crit calculation"
##   factor_df residual_df   F_crit
## 1         1          596 3.857108
```


## Other plots

The following are density plots for experiments ran between 02/25/2024 and 03/20/2024.


On the first density plot, it is limited between 0 and 1. Its hiding a mode for Lambda\_DD on Tier 1 that happens after 25 and 50 ms.

Density plots


Write latency for a 10 KB file and 512 B I/O size




Write latency for a 1 GB file and 512 B I/O size




# Write latency for a 1 GB file and 128 KB I/O size




Read latency for a 10 KB file and 512 B I/O size










Read latency for a 1 GB file and 128 KB I/O size




Box plots




Write latency for a 1 GB file and 128 KB I/O size





Read latency for a 1 GB file and 512 B I/O size



Read latency for a 1 GB file and 128 KB I/O size

